

Race conditions

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Programming Concepts using Java

Week 10

Threads and shared variables

- Threads are lightweight processes with shared variables that can run in parallel
- Browser example: download thread and user-interface thread run in parallel
 - Shared boolean variable `terminate` indicates whether download should be interrupted
 - `terminate` is initially false
 - Clicking `Stop` sets it to true
 - Download thread checks the value of this variable periodically and aborts if it is set to true
- Watch out for **race conditions**
 - Shared variables must be updated consistently

Maintaining data consistency

- `double accounts[100]` describes 100 bank accounts

Maintaining data consistency

- `double accounts[100]` describes 100 bank accounts
- Two functions that operate on `accounts`: `transfer()` and `audit()`

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}
```

```
double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Maintaining data consistency

- `double accounts[100]` describes 100 bank accounts
- Two functions that operate on `accounts`: `transfer()` and `audit()`
- What are the possibilities when we execute the following?

Thread 1

```
...  
status =  
    transfer(500.00,7,8);  
...
```

Thread 2

```
...  
System.out.  
    print(audit());  
...
```

```
boolean transfer (double amount,  
                  int source,  
                  int target){  
    if (accounts[source] < amount){  
        return false;  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}  
  
double audit(){  
    // total balance across all accounts  
    double balance = 0.00;  
    for (int i = 0; i < 100; i++){  
        balance += accounts[i];  
    }  
    return balance;  
}
```

Maintaining data consistency ...

- What are the possibilities when we execute the following?

Thread 1

...

status =

transfer(500.00,7,8);

...

Thread 2

...

System.out.

print(audit());

...

- `audit()` can report an overall total that is 500 more or less than the actual assets

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}
```

```
double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Maintaining data consistency ...

- What are the possibilities when we execute the following?

Thread 1

...

status =

transfer(500.00,7,8);

...

Thread 2

...

System.out.

print(audit());

...

- `audit()` can report an overall total that is 500 more or less than the actual assets

- Depends on how actions of `transfer` are interleaved with actions of `audit`

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}
```

```
double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Maintaining data consistency ...

- What are the possibilities when we execute the following?

Thread 1

...

status =

transfer(500.00,7,8);

...

Thread 2

...

System.out.

print(audit());

...

- `audit()` can report an overall total that is 500 more or less than the actual assets

- Depends on how actions of `transfer` are interleaved with actions of `audit`
- Can even report an error if `transfer` happens **atomically**

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}
```

```
double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```


Atomicity of updates

- Two threads increment a shared variable `n`

Thread 1

```
...  
m = n;  
m++;  
n = m;  
...
```

Thread 2

```
...  
k = n;  
k++;  
n = k;  
...
```

Atomicity of updates

- Two threads increment a shared variable `n`

Thread 1

```
...  
m = n;  
m++;  
n = m;  
...
```

Thread 2

```
...  
k = n;  
k++;  
n = k;  
...
```

- Expect `n` to increase by 2 ...

Atomicity of updates

- Two threads increment a shared variable `n`

Thread 1	Thread 2
...	...
<code>m = n;</code>	<code>k = n;</code>
<code>m++;</code>	<code>k++;</code>
<code>n = m;</code>	<code>n = k;</code>
...	...

- Expect `n` to increase by 2 ...
- ... but, time-slicing may order execution as follows

```
Thread 1: m = n;  
Thread 1: m++;  
Thread 2: k = n;    // k gets the original value of n  
Thread 2: k++;  
Thread 1: n = m;  
Thread 2: n = k;    // Same value as that set by Thread 1
```

Race conditions and mutual exclusion

- **Race condition** — concurrent update of shared variables, unpredictable outcome
 - Executing `transfer()` and `audit()` concurrently can cause `audit()` to report more or less than the actual assets

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}

double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Race conditions and mutual exclusion

- **Race condition** — concurrent update of shared variables, unpredictable outcome
 - Executing `transfer()` and `audit()` concurrently can cause `audit()` to report more or less than the actual assets
- Avoid this by insisting that `transfer()` and `audit()` do not interleave

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}
```

```
double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Race conditions and mutual exclusion

- **Race condition** — concurrent update of shared variables, unpredictable outcome
 - Executing `transfer()` and `audit()` concurrently can cause `audit()` to report more or less than the actual assets
- Avoid this by insisting that `transfer()` and `audit()` do not interleave
- Never simultaneously have current control point of one thread within `transfer()` and another thread within `audit()`

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}

double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Race conditions and mutual exclusion

- **Race condition** — concurrent update of shared variables, unpredictable outcome
 - Executing `transfer()` and `audit()` concurrently can cause `audit()` to report more or less than the actual assets
- Avoid this by insisting that `transfer()` and `audit()` do not interleave
- Never simultaneously have current control point of one thread within `transfer()` and another thread within `audit()`
- **Mutually exclusive** access to **critical regions** of code

```
boolean transfer (double amount,
                  int source,
                  int target){
    if (accounts[source] < amount){
        return false;
    }
    accounts[source] -= amount;
    accounts[target] += amount;
    return true;
}
```

```
double audit(){
    // total balance across all accounts
    double balance = 0.00;
    for (int i = 0; i < 100; i++){
        balance += accounts[i];
    }
    return balance;
}
```

Summary

- Concurrent update of a shared variable can lead to data inconsistency
 - Race condition
- Control behaviour of threads to regulate concurrent updates
 - Critical sections — sections of code where shared variables are updated
 - Mutual exclusion — at most one thread at a time can be in a critical section