

# Monitors in Java

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Programming Concepts using Java

Week 11

# Monitors

- Monitor is like a class in an OO language
  - Data definition — to which access is restricted across threads
  - Collections of functions operating on this data — all are implicitly mutually exclusive
- Monitor guarantees mutual exclusion — if one function is active, any other function will have to wait for it to finish
- Implicit **queue** associated with each monitor
  - Contains all processes waiting for access

```
monitor bank_account{
    double accounts[100];

    boolean transfer (double amount,
                     int source,
                     int target){
        if (accounts[source] < amount){
            return false;
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        return true;
    }

    double audit(){
        // compute balance across all accounts
        double balance = 0.00;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }
}
```

# Condition variables

- Thread suspends itself and waits for a state change — `q[source].wait()`
- Separate **internal** queue, vs **external** queue for initially blocked threads

```
monitor bank_account{
    double accounts[100];
    queue q[100]; // one internal queue
                  // for each account
    boolean transfer (double amount,
                     int source,
                     int target){
        while (accounts[source] < amount){
            q[source].wait(); // wait in the queue
                              // associated with source
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        q[target].notify(); // notify the queue
                            // associated with target
        return true;
    }

    // compute the balance across all accounts
    double audit(){ ...}
}
```

# Condition variables

- Thread suspends itself and waits for a state change — `q[source].wait()`
- Separate **internal** queue, vs **external** queue for initially blocked threads
- Notify change — `q[target].notify()`

```
monitor bank_account{
    double accounts[100];
    queue q[100]; // one internal queue
                  // for each account
    boolean transfer (double amount,
                      int source,
                      int target){
        while (accounts[source] < amount){
            q[source].wait(); // wait in the queue
                              // associated with source
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        q[target].notify(); // notify the queue
                             // associated with target
        return true;
    }

    // compute the balance across all accounts
    double audit(){ ...}
}
```

# Condition variables

- Thread suspends itself and waits for a state change — `q[source].wait()`
- Separate **internal** queue, vs **external** queue for initially blocked threads
- Notify change — `q[target].notify()`
- **Signal and exit** — notifying process immediately exits the monitor
- **Signal and wait** — notifying process swaps roles with notified process
- **Signal and continue** — notifying process keeps control till it completes and then one of the notified processes steps in

```
monitor bank_account{
    double accounts[100];
    queue q[100]; // one internal queue
                  // for each account
    boolean transfer (double amount,
                     int source,
                     int target){
        while (accounts[source] < amount){
            q[source].wait(); // wait in the queue
                              // associated with source
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        q[target].notify(); // notify the queue
                             // associated with target
        return true;
    }

    // compute the balance across all accounts
    double audit(){ ...}
}
```

# Monitors in Java

- Monitors incorporated within existing class definitions

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```

# Monitors in Java

- Monitors incorporated within existing class definitions
- Function declared `synchronized` is to be executed atomically

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```

# Monitors in Java

- Monitors incorporated within existing class definitions
- Function declared `synchronized` is to be executed atomically
- Each object has a `lock`
  - To execute a `synchronized` method, thread must acquire lock
  - Thread gives up lock when the method exits
  - Only one thread can have the lock at any time

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```



# Monitors in Java

- Monitors incorporated within existing class definitions
- Function declared `synchronized` is to be executed atomically
- Each object has a `lock`
  - To execute a `synchronized` method, thread must acquire lock
  - Thread gives up lock when the method exits
  - Only one thread can have the lock at any time
- Wait for lock in external queue

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```

# Monitors in Java

- `wait()` and `notify()` to suspend and resume

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```

# Monitors in Java

- `wait()` and `notify()` to suspend and resume
- Wait — single internal queue

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```

# Monitors in Java

- `wait()` and `notify()` to suspend and resume
- Wait — single internal queue
- Notify
  - `notify()` signals one (arbitrary) waiting process
  - `notifyAll()` signals all waiting processes
  - Java uses **signal and continue**

```
public class bank_account{
    double accounts[100];

    public synchronized boolean
        transfer(double amount, int source, int target){
        while (accounts[source] < amount){ wait(); }
        accounts[source] -= amount;
        accounts[target] += amount;
        notifyAll();
        return true;
    }

    public synchronized double audit(){
        double balance = 0.0;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }

    public double current_balance(int i){
        return accounts[i];    // not synchronized!
    }
}
```

# Object locks ...

- Use object locks to synchronize arbitrary blocks of code

```
public class XYZ{  
    Object o = new Object();  
  
    public int f(){  
        ..  
        synchronized(o){ ... }  
    }  
  
    public double g(){  
        ..  
        synchronized(o){ ... }  
    }  
}
```

# Object locks ...

- Use object locks to synchronize arbitrary blocks of code
- `f()` and `g()` can start in parallel
- Only one of the threads can grab the lock for `o`

```
public class XYZ{  
    Object o = new Object();  
  
    public int f(){  
        ..  
        synchronized(o){ ... }  
    }  
  
    public double g(){  
        ..  
        synchronized(o){ ... }  
    }  
}
```

# Object locks ...

- Use object locks to synchronize arbitrary blocks of code
- `f()` and `g()` can start in parallel
- Only one of the threads can grab the lock for `o`
- Each object has its own internal queue

```
Object o = new Object();

public int f(){
    ..
    synchronized(o){
        ...
        o.wait();    // Wait in queue attached to "o"
        ...
    }
}

public double g(){
    ..
    synchronized(o){
        ...
        o.notifyAll();    // Wake up queue attached to
        ...
    }
}
```

# Object locks ...

- Use object locks to synchronize arbitrary blocks of code
- `f()` and `g()` can start in parallel
- Only one of the threads can grab the lock for `o`
- Each object has its own internal queue
- Can convert methods from “externally” synchronized to “internally” synchronized

```
public double h(){  
    synchronized(this){  
        ...  
    }  
}
```



# Object locks ...

- Use object locks to synchronize arbitrary blocks of code
- `f()` and `g()` can start in parallel
- Only one of the threads can grab the lock for `o`
- Each object has its own internal queue
- Can convert methods from “externally” synchronized to “internally” synchronized
- “Anonymous” `wait()`, `notify()`, `notifyAll()` abbreviate `this.wait()`, `this.notify()`, `this.notifyAll()`

```
public double h(){  
    synchronized(this){  
        ...  
    }  
}
```

## Object locks ...

- Actually, `wait()` can be “interrupted” by an `InterruptedException`

# Object locks ...

- Actually, `wait()` can be “interrupted” by an `InterruptedException`
- Should write

```
try{  
    wait();  
}  
catch (InterruptedException e) {  
    ...  
};
```

# Object locks ...

- Actually, `wait()` can be “interrupted” by an `InterruptedException`

- Should write

```
try{
    wait();
}
catch (InterruptedException e) {
    ...
};
```

- Error to use `wait()`, `notify()`, `notifyAll()` outside synchronized method
  - `IllegalMonitorStateException`

# Object locks ...

- Actually, `wait()` can be “interrupted” by an `InterruptedException`

- Should write

```
try{
    wait();
}
catch (InterruptedException e) {
    ...
};
```

- Error to use `wait()`, `notify()`, `notifyAll()` outside synchronized method

- `IllegalMonitorStateException`

- Likewise, use `o.wait()`, `o.notify()`, `o.notifyAll()` only in block synchronized on `o`

# Reentrant locks

- Separate `ReentrantLock` class

```
public class Bank
{
    private Lock bankLock = new ReentrantLock();
    ...
    public void
        transfer(int from, int to, int amount) {
        bankLock.lock();
        try {
            accounts[from] -= amount;
            accounts[to] += amount;
        }
        finally {
            bankLock.unlock();
        }
    }
}
```

# Reentrant locks

- Separate `ReentrantLock` class
- Similar to a semaphore
  - `lock()` is like  $P(S)$
  - `unlock()` is like  $V(S)$

```
public class Bank
{
    private Lock bankLock = new ReentrantLock();
    ...
    public void
        transfer(int from, int to, int amount) {
        bankLock.lock();
        try {
            accounts[from] -= amount;
            accounts[to] += amount;
        }
        finally {
            bankLock.unlock();
        }
    }
}
```

# Reentrant locks

- Separate `ReentrantLock` class
- Similar to a semaphore
  - `lock()` is like  $P(S)$
  - `unlock()` is like  $V(S)$
- Always `unlock()` in `finally` — avoid abort while holding lock

```
public class Bank
{
    private Lock bankLock = new ReentrantLock();
    ...
    public void
        transfer(int from, int to, int amount) {
        bankLock.lock();
        try {
            accounts[from] -= amount;
            accounts[to] += amount;
        }
        finally {
            bankLock.unlock();
        }
    }
}
```



# Reentrant locks

- Separate `ReentrantLock` class
- Similar to a semaphore
  - `lock()` is like  $P(S)$
  - `unlock()` is like  $V(S)$
- Always `unlock()` in `finally` — avoid abort while holding lock
- Why **reentrant**?
  - Thread holding lock can reacquire it
  - `transfer()` may call `getBalance()` that also locks `bankLock`
  - **Hold count** increases with `lock()`, decreases with `unlock()`
  - Lock is available if hold count is 0

```
public class Bank
{
    private Lock bankLock = new ReentrantLock();
    ...
    public void
        transfer(int from, int to, int amount) {
        bankLock.lock();
        try {
            accounts[from] -= amount;
            accounts[to] += amount;
        }
        finally {
            bankLock.unlock();
        }
    }
}
```

# Summary

- Every object in Java implicitly has a lock
- Methods tagged `synchronized` are executed atomically
  - Implicitly acquire and release the object's lock
- Associated condition variable, single internal queue
  - `wait()`, `notify()`, `notifyAll()`
- Can synchronize an arbitrary block of code using an object
  - `synchronized(o) { ... }`
  - `o.wait()`, `o.notify()`, `o.notifyAll()`
- Reentrant locks work like semaphores