

# Mutual Exclusion

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Programming Concepts using Java

Week 10

# Mutual exclusion

- Concurrent update of a shared variable can lead to data inconsistency
  - Race condition
- Control behaviour of threads to regulate concurrent updates
  - Critical sections — sections of code where shared variables are updated
  - Mutual exclusion — at most one thread at a time can be in a critical section

# Mutual exclusion for two processes

## ■ First attempt

Thread 1

```
...  
while (turn != 1){  
    // "Busy" wait  
}  
// Enter critical section  
...  
// Leave critical section  
turn = 2;  
...
```

Thread 2

```
...  
while (turn != 2){  
    // "Busy" wait  
}  
// Enter critical section  
...  
// Leave critical section  
turn = 1;  
...
```

# Mutual exclusion for two processes

## ■ First attempt

Thread 1

```
...  
while (turn != 1){  
    // "Busy" wait  
}  
// Enter critical section  
...  
// Leave critical section  
turn = 2;  
...
```

Thread 2

```
...  
while (turn != 2){  
    // "Busy" wait  
}  
// Enter critical section  
...  
// Leave critical section  
turn = 1;  
...
```

- Shared variable `turn` — no assumption about initial value, atomic update

# Mutual exclusion for two processes

## ■ First attempt

Thread 1

```
...  
while (turn != 1){  
    // "Busy" wait  
}  
// Enter critical section  
...  
// Leave critical section  
turn = 2;  
...
```

Thread 2

```
...  
while (turn != 2){  
    // "Busy" wait  
}  
// Enter critical section  
...  
// Leave critical section  
turn = 1;  
...
```

- Shared variable `turn` — no assumption about initial value, atomic update
- Mutually exclusive access is guaranteed ...

# Mutual exclusion for two processes

## ■ First attempt

```
Thread 1
...
while (turn != 1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
turn = 2;
...
```

```
Thread 2
...
while (turn != 2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
turn = 1;
...
```

- Shared variable `turn` — no assumption about initial value, atomic update
- Mutually exclusive access is guaranteed ...
- ...but one thread is locked out permanently if other thread shuts down
  - **Starvation!**

# Mutual exclusion for two processes ...

## ■ Second attempt

Thread 1

```
...
request_1 = true;
while (request_2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_1 = false;
...
```

Thread 2

```
...
request_2 = true;
while (request_1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_2 = false;
...
```

# Mutual exclusion for two processes ...

## ■ Second attempt

Thread 1

```
...
request_1 = true;
while (request_2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_1 = false;
...
```

Thread 2

```
...
request_2 = true;
while (request_1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_2 = false;
...
```

## ■ Mutually exclusive access is guaranteed ...



# Mutual exclusion for two processes ...

## ■ Second attempt

```
Thread 1
...
request_1 = true;
while (request_2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_1 = false;
...
```

```
Thread 2
...
request_2 = true;
while (request_1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_2 = false;
...
```

## ■ Mutually exclusive access is guaranteed ...

## ■ ...but if both threads try simultaneously, they block each other

### ■ **Deadlock!**

# Peterson's algorithm

Thread 1

```
...
request_1 = true;
turn = 2;
while (request_2 &&
       turn != 1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_1 = false;
...
```

Thread 2

```
...
request_2 = true;
turn = 1;
while (request_1 &&
       turn != 2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_2 = false;
...
```

- Combines the previous two approaches

# Peterson's algorithm

```
Thread 1
...
request_1 = true;
turn = 2;
while (request_2 &&
      turn != 1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_1 = false;
...
```

```
Thread 2
...
request_2 = true;
turn = 1;
while (request_1 &&
      turn != 2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_2 = false;
...
```

- Combines the previous two approaches
- If both try simultaneously, **turn** decides who goes through

# Peterson's algorithm

```
Thread 1
...
request_1 = true;
turn = 2;
while (request_2 &&
       turn != 1){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_1 = false;
...
```

```
Thread 2
...
request_2 = true;
turn = 1;
while (request_1 &&
       turn != 2){
    // "Busy" wait
}
// Enter critical section
...
// Leave critical section
request_2 = false;
...
```

- Combines the previous two approaches
- If both try simultaneously, **turn** decides who goes through
- If only one is alive, **request** for that process is stuck at false and **turn** is irrelevant

# Beyond two processes

- Generalizing Peterson's solution to more than two processes is not trivial

# Beyond two processes

- Generalizing Peterson's solution to more than two processes is not trivial
- For  $n$  process mutual exclusion other solutions exist

# Beyond two processes

- Generalizing Peterson's solution to more than two processes is not trivial
- For  $n$  process mutual exclusion other solutions exist
- Lamport's **Bakery Algorithm**
  - Each new process picks up a token (increments a counter) that is larger than all waiting processes
  - Lowest token number gets served next
  - Still need to break ties — token counter is not atomic

# Beyond two processes

- Generalizing Peterson's solution to more than two processes is not trivial
- For  $n$  process mutual exclusion other solutions exist
- Lamport's **Bakery Algorithm**
  - Each new process picks up a token (increments a counter) that is larger than all waiting processes
  - Lowest token number gets served next
  - Still need to break ties — token counter is not atomic
- Need specific clever solutions for different situations



# Beyond two processes

- Generalizing Peterson's solution to more than two processes is not trivial
- For  $n$  process mutual exclusion other solutions exist
- Lamport's **Bakery Algorithm**
  - Each new process picks up a token (increments a counter) that is larger than all waiting processes
  - Lowest token number gets served next
  - Still need to break ties — token counter is not atomic
- Need specific clever solutions for different situations
- Need to argue correctness in each case

# Beyond two processes

- Generalizing Peterson's solution to more than two processes is not trivial
- For  $n$  process mutual exclusion other solutions exist
- Lamport's **Bakery Algorithm**
  - Each new process picks up a token (increments a counter) that is larger than all waiting processes
  - Lowest token number gets served next
  - Still need to break ties — token counter is not atomic
- Need specific clever solutions for different situations
- Need to argue correctness in each case
- Instead, provide higher level support in programming language for synchronization

# Summary

- We can construct protocols that guarantee mutual exclusion to critical sections
  - Watch out for **starvation** and **deadlock**
- These protocols cleverly use regular variables
  - No assumptions about initial values, atomicity of updates
- Difficult to generalize such protocols to arbitrary situations
- Look to programming language for features that control synchronization