

# Monitors

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Programming Concepts using Java

Week 10

# Atomic test-and-set

- Test-and-set is at the heart of most race conditions
- Need a high level primitive for atomic test-and-set in the programming language
- Semaphores provide one such solution
- Solutions based on test-and-set are low level and prone to programming errors

# Monitors

- Attach synchronization control to the data that is being protected

# Monitors

- Attach synchronization control to the data that is being protected
- **Monitors** — Per Brinch Hansen and CAR Hoare

# Monitors

- Attach synchronization control to the data that is being protected
- **Monitors** — Per Brinch Hansen and CAR Hoare
- Monitor is like a class in an OO language
  - Data definition — to which access is restricted across threads
  - Collections of functions operating on this data — all are implicitly mutually exclusive

```
monitor bank_account{
    double accounts[100];

    boolean transfer (double amount,
                     int source,
                     int target){
        if (accounts[source] < amount){
            return false;
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        return true;
    }

    double audit(){
        // compute balance across all accounts
        double balance = 0.00;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }
}
```

# Monitors

- Attach synchronization control to the data that is being protected
- **Monitors** — Per Brinch Hansen and CAR Hoare
- Monitor is like a class in an OO language
  - Data definition — to which access is restricted across threads
  - Collections of functions operating on this data — all are implicitly mutually exclusive
- Monitor guarantees mutual exclusion — if one function is active, any other function will have to wait for it to finish

```
monitor bank_account{
    double accounts[100];

    boolean transfer (double amount,
                      int source,
                      int target){
        if (accounts[source] < amount){
            return false;
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        return true;
    }

    double audit(){
        // compute balance across all accounts
        double balance = 0.00;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }
}
```

# Monitors: external queue

- Monitor ensures **transfer** and **audit** are mutually exclusive

```
monitor bank_account{
    double accounts[100];

    boolean transfer (double amount,
                      int source,
                      int target){
        if (accounts[source] < amount){
            return false;
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        return true;
    }

    double audit(){
        // compute balance across all accounts
        double balance = 0.00;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }
}
```

# Monitors: external queue

- Monitor ensures **transfer** and **audit** are mutually exclusive
- If **Thread 1** is executing **transfer** and **Thread 2** invokes **audit**, it must wait

```
monitor bank_account{
    double accounts[100];

    boolean transfer (double amount,
                      int source,
                      int target){
        if (accounts[source] < amount){
            return false;
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        return true;
    }

    double audit(){
        // compute balance across all accounts
        double balance = 0.00;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }
}
```



# Monitors: external queue

- Monitor ensures **transfer** and **audit** are mutually exclusive
- If **Thread 1** is executing **transfer** and **Thread 2** invokes **audit**, it must wait
- Implicit **queue** associated with each monitor
  - Contains all processes waiting for access
  - In practice, this may be just a set, not a queue

```
monitor bank_account{
    double accounts[100];

    boolean transfer (double amount,
                      int source,
                      int target){
        if (accounts[source] < amount){
            return false;
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        return true;
    }

    double audit(){
        // compute balance across all accounts
        double balance = 0.00;
        for (int i = 0; i < 100; i++){
            balance += accounts[i];
        }
        return balance;
    }
}
```

# Making monitors more flexible

- Our definition of monitors may be too restrictive

```
transfer(500.00,i,j);
```

```
transfer(400.00,j,k);
```

# Making monitors more flexible

- Our definition of monitors may be too restrictive

```
transfer(500.00,i,j);  
transfer(400.00,j,k);
```

- This should always succeed if `accounts[i] > 500`

# Making monitors more flexible

- Our definition of monitors may be too restrictive

```
transfer(500.00,i,j);  
transfer(400.00,j,k);
```

- This should always succeed if `accounts[i] > 500`
- If these calls are reordered and `accounts[j] < 400` initially, this will fail

# Making monitors more flexible

- Our definition of monitors may be too restrictive

```
transfer(500.00,i,j);  
transfer(400.00,j,k);
```

- This should always succeed if `accounts[i] > 500`
- If these calls are reordered and `accounts[j] < 400` initially, this will fail
- A possible fix — let an account wait for pending inflows

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){  
        // wait for another transaction to transfer money  
        // into accounts[source]  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}
```

# Monitors — wait()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){  
        // wait for another transaction to transfer money  
        // into accounts[source]  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}
```

- All other processes are blocked out while this process waits!

# Monitors — wait()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){  
        // wait for another transaction to transfer money  
        // into accounts[source]  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}
```

- All other processes are blocked out while this process waits!
- Need a mechanism for a thread to suspend itself and give up the monitor

# Monitors — wait()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){  
        // wait for another transaction to transfer money  
        // into accounts[source]  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}
```

- All other processes are blocked out while this process waits!
- Need a mechanism for a thread to suspend itself and give up the monitor
- A suspended process is waiting for monitor to change its state



# Monitors — wait()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){  
        // wait for another transaction to transfer money  
        // into accounts[source]  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}
```

- All other processes are blocked out while this process waits!
- Need a mechanism for a thread to suspend itself and give up the monitor
- A suspended process is waiting for monitor to change its state
- Have a separate **internal** queue, as opposed to **external** queue where initially blocked threads wait

# Monitors — wait()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){  
        // wait for another transaction to transfer money  
        // into accounts[source]  
    }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    return true;  
}
```

- All other processes are blocked out while this process waits!
- Need a mechanism for a thread to suspend itself and give up the monitor
- A suspended process is waiting for monitor to change its state
- Have a separate **internal** queue, as opposed to **external** queue where initially blocked threads wait
- Dual operation to **notify** and wake up suspended processes

# Monitors — notify()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){ wait(); }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    notify();  
    return true;  
}
```

# Monitors — notify()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){ wait(); }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    notify();  
    return true;  
}
```

- What happens when a process executes `notify()`?

# Monitors — notify()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){ wait(); }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    notify();  
    return true;  
}
```

- What happens when a process executes `notify()`?
- **Signal and exit** — notifying process immediately exits the monitor
  - `notify()` must be the last instruction

# Monitors — notify()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){ wait(); }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    notify();  
    return true;  
}
```

- What happens when a process executes `notify()`?
- **Signal and exit** — notifying process immediately exits the monitor
  - `notify()` must be the last instruction
- **Signal and wait** — notifying process swaps roles and goes into the internal queue of the monitor

# Monitors — notify()

```
boolean transfer (double amount, int source, int target){  
    if (accounts[source] < amount){ wait(); }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    notify();  
    return true;  
}
```

- What happens when a process executes `notify()`?
- **Signal and exit** — notifying process immediately exits the monitor
  - `notify()` must be the last instruction
- **Signal and wait** — notifying process swaps roles and goes into the internal queue of the monitor
- **Signal and continue** — notifying process keeps control till it completes and then one of the notified processes steps in

# Monitors — `wait()` and `notify()`

- Should check the `wait()` condition again on wake up
  - Change of state may not be sufficient to continue — e.g., not enough inflow into the account to allow transfer



# Monitors — `wait()` and `notify()`

- Should check the `wait()` condition again on wake up
  - Change of state may not be sufficient to continue — e.g., not enough inflow into the account to allow transfer
- A thread can be again interleaved between notification and running
  - At wake-up, the state was fine, but it has changed again due to some other concurrent action

# Monitors — wait() and notify()

- Should check the `wait()` condition again on wake up
  - Change of state may not be sufficient to continue — e.g., not enough inflow into the account to allow transfer
- A thread can be again interleaved between notification and running
  - At wake-up, the state was fine, but it has changed again due to some other concurrent action
- `wait()` should be in a `while`, not in an `if`

```
boolean transfer (double amount, int source, int target){  
    while (accounts[source] < amount){ wait(); }  
    accounts[source] -= amount;  
    accounts[target] += amount;  
    notify();  
    return true;  
}
```

# Condition variables

- After `transfer`, `notify()` is only useful for threads waiting for target account of transfer to change state

# Condition variables

- After `transfer`, `notify()` is only useful for threads waiting for target account of transfer to change state
- Makes sense to have more than one internal queue

# Condition variables

- After `transfer`, `notify()` is only useful for threads waiting for target account of transfer to change state
- Makes sense to have more than one internal queue
- Monitor can have **condition variables** to describe internal queues

```
monitor bank_account{
    double accounts[100];
    queue q[100]; // one internal queue
                // for each account
    boolean transfer (double amount,
                     int source,
                     int target){
        while (accounts[source] < amount){
            q[source].wait(); // wait in the queue
                             // associated with source
        }
        accounts[source] -= amount;
        accounts[target] += amount;
        q[target].notify(); // notify the queue
                           // associated with target
        return true;
    }

    // compute the balance across all accounts
    double audit(){ ...}
}
```

# Summary

- Concurrent programming with atomic test-and-set primitives is error prone
- Monitors are like abstract datatypes for concurrent programming
  - Encapsulate data and methods to manipulate data
  - Methods are implicitly atomic, regulate concurrent access
  - Each object has an implicit external queue of processes waiting to execute a method
- `wait()` and `notify()` allow more flexible operation
- Can have multiple internal queues controlled by condition variables