

Single Source Shortest Paths

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Programming, Data Structures and Algorithms using Python
Week 5

Single source shortest paths


- Weighted graph:

- $G = (V, E)$

- $W : E \rightarrow \mathbb{R}$

- Single source shortest paths

- Find shortest paths from a fixed vertex to every other vertex



Single source shortest paths

- Weighted graph:


- $G = (V, E)$

- $W : E \rightarrow \mathbb{R}$

- Single source shortest paths


- Find shortest paths from a fixed vertex to every other vertex

- Assume, for now, that edge weights are all non-negative




Single source shortest paths

- Compute shortest paths from 0 to all other vertices




Single source shortest paths

- Compute shortest paths from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines




Single source shortest paths

- Compute shortest paths from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0




Single source shortest paths

- Compute shortest paths from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline




Single source shortest paths

- Compute shortest paths from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex




Single source shortest paths

- Compute shortest paths from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex




Single source shortest paths

- Compute shortest paths from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours




Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours



Single source shortest paths

- Compute **expected burn time** for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours
- Algorithm due to **Edsger W Dijkstra**



Dijkstra's algorithm: Proof of correctness


- Each new shortest path we discover extends an earlier one

Dijkstra's algorithm: Proof of correctness

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt


Dijkstra's algorithm: Proof of correctness

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt




Dijkstra's algorithm: Proof of correctness

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is **v**, via **x**




Dijkstra's algorithm: Proof of correctness

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is **v**, via **x**
- Cannot find a shorter path later from **y** to **v** via **w**
 - Burn time of **w** \geq burn time of **v**
 - Edge from **w** to **v** has weight ≥ 0



Dijkstra's algorithm: Proof of correctness

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is v , via x
- Cannot find a shorter path later from y to v via w
 - Burn time of $w \geq$ burn time of v
 - Edge from w to v has weight ≥ 0
- This argument breaks down if edge (w,v) can have negative weight
 - Can't use Dijkstra's algorithm with negative edge weights



Implementation

- Maintain two dictionaries with vertices as keys
 - `visited`, initially `False` for all `v` (burnt vertices)
 - `distance`, initially `infinity` for all `v` (expected burn time)

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Implementation

- Maintain two dictionaries with vertices as keys
 - `visited`, initially `False` for all `v` (burnt vertices)
 - `distance`, initially `infinity` for all `v` (expected burn time)
- Set `distance[s]` to 0

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```


Implementation

- Maintain two dictionaries with vertices as keys
 - `visited`, initially `False` for all `v` (burnt vertices)
 - `distance`, initially `infinity` for all `v` (expected burn time)
- Set `distance[s]` to 0
- Repeat, until all reachable vertices are visited
 - Find unvisited vertex `nextv` with minimum distance
 - Set `visited[nextv]` to `True`
 - Recompute `distance[v]` for every neighbour `v` of `nextv`

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Complexity

- Setting `infinity` takes $O(n^2)$ time

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Complexity

- Setting `infinity` takes $O(n^2)$ time
- Main loop runs n times
 - Each iteration visits one more vertex
 - $O(n)$ to find next vertex to visit
 - $O(n)$ to update `distance[v]` for neighbours

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Complexity

- Setting `infinity` takes $O(n^2)$ time
- Main loop runs n times
 - Each iteration visits one more vertex
 - $O(n)$ to find next vertex to visit
 - $O(n)$ to update `distance[v]` for neighbours
- Overall $O(n^2)$

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Complexity

- Setting `infinity` takes $O(n^2)$ time
- Main loop runs n times
 - Each iteration visits one more vertex
 - $O(n)$ to find next vertex to visit
 - $O(n)$ to update `distance[v]` for neighbours
- Overall $O(n^2)$
- If we use an adjacency list
 - Setting `infinity` and updating distances both $O(m)$, amortised
 - $O(n)$ bottleneck remains to find next vertex to visit
 - Better data structure? Later ...

```
def dijkstralist(WList,s):
    infinity = 1 + len(WList.keys())*
                max([d for u in WList.keys()
                     for (v,d) in WList[u]])
    (visited,distance) = ({},{})
    for v in WList.keys():
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in WList.keys():
        nextd = min([distance[v] for v in WList.keys()
                     if not visited[v]])
        nextvlist = [v for v in WList.keys()
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for (v,d) in WList[nextv]:
            if not visited[v]:
                distance[v] = min(distance[v],distance[nextv]+d)
    return(distance)
```

Summary

- Dijkstra's algorithm computes single source shortest paths
- Use a **greedy** strategy to identify vertices to visit
 - Next vertex to visit is based on shortest distance computed so far
 - Need to prove that such a strategy is correct
 - Correctness requires edge weights to be non-negative
- Complexity is $O(n^2)$
 - Even with adjacency lists
 - Bottleneck is identifying unvisited vertex with minimum distance
 - Need a better data structure to identify and remove minimum (or maximum) from a collection