

Using Heaps in Algorithms

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Programming, Data Structures and Algorithms using Python

Week 6

Priority queues and heaps

- Priority queues support the following operations
 - `insert()`
 - `delete_max()` or `delete_min()`
- Heaps are a tree based implementation of priority queues
 - `insert()`, `delete_max()` / `delete_min()` are both $O(\log n)$
 - `heapify()` builds a heap from a list/array in time $O(n)$
- Heap can be represented as a list/array
 - Simple index arithmetic to find parent and children of a node
- What more do we need to use a heap in an algorithm?

Dijkstra's algorithm

- Maintain two dictionaries with vertices as keys
 - `visited`, initially `False` for all `v`
 - `distance`, initially `infinity` for all `v`
- Set `distance[s]` to 0
- Repeat, until all reachable vertices are visited
 - Find unvisited vertex `nextv` with minimum distance
 - Set `visited[nextv]` to `True`
 - Recompute `distance[v]` for every neighbour `v` of `nextv`

```
def dijkstra(WMat,s):  
    (rows,cols,x) = WMat.shape  
    infinity = np.max(WMat)*rows+1  
    (visited,distance) = ({},{})  
    for v in range(rows):  
        (visited[v],distance[v]) = (False,infinity)  
    distance[s] = 0  
    for u in range(rows):  
        nextd = min([distance[v] for v in range(rows)  
                     if not visited[v]])  
        nextvlist = [v for v in range(rows)  
                     if (not visited[v]) and  
                        distance[v] == nextd]  
  
        if nextvlist == []:  
            break  
        nextv = min(nextvlist)  
        visited[nextv] = True  
        for v in range(cols):  
            if WMat[nextv,v,0] == 1 and (not visited[v]):  
                distance[v] = min(distance[v],distance[nextv]  
                                   +WMat[nextv,v,1])  
  
    return(distance)
```

Dijkstra's algorithm

Bottleneck

- Find unvisited vertex j with minimum distance
 - Naive implementation requires an $O(n)$ scan

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Dijkstra's algorithm

Bottleneck

- Find unvisited vertex j with minimum distance
 - Naive implementation requires an $O(n)$ scan
- Maintain unvisited vertices as a min-heap
 - `delete_min()` in $O(\log n)$ time

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```

Dijkstra's algorithm


Bottleneck

- Find unvisited vertex j with minimum distance
 - Naive implementation requires an $O(n)$ scan
- Maintain unvisited vertices as a min-heap
 - `delete_min()` in $O(\log n)$ time
- But, also need to update distances of neighbours
 - Unvisited neighbours' distances are inside the min-heap
 - Updating a value is not a basic heap operation

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})
    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
                distance[v] = min(distance[v],distance[nextv]
                                   +WMat[nextv,v,1])
    return(distance)
```


Updating values in a min-heap

- Change 54 to 35




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`
- Change 29 to 71




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`
- Change 29 to 71
 - Increasing a value can create a violation with child




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`
- Change 29 to 71
 - Increasing a value can create a violation with child
 - Swap downwards to restore heap, similar to `delete_min()`




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`
- Change 29 to 71
 - Increasing a value can create a violation with child
 - Swap downwards to restore heap, similar to `delete_min()`




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`
- Change 29 to 71
 - Increasing a value can create a violation with child
 - Swap downwards to restore heap, similar to `delete_min()`
- Both updates are $O(\log n)$
 - Are we done?




Updating values in a min-heap

- Change 54 to 35
 - Reducing a value can create a violation with parent
 - Swap upwards to restore heap, similar to `insert()`
- Change 29 to 71
 - Increasing a value can create a violation with child
 - Swap downwards to restore heap, similar to `delete_min()`
- Both updates are $O(\log n)$
 - Are we done?
- Locate the node to update?



Updating values in a min-heap

- Maintain two additional dictionaries
 - Vertices are $\{0, 1, \dots, n-1\}$
 - Heap positions are $\{0, 1, \dots, n-1\}$
 - **VtoH** maps vertices to heap positions
 - **HtoV** maps heap positions to vertices



VtoH

0	1	2	3	4	5	6	7	8
7	3	5	2	4	1	6	0	8

HtoV


0	1	2	3	4	5	6	7	8
7	5	3	1	4	2	6	0	8

Updating values in a min-heap

- Maintain two additional dictionaries

- Vertices are $\{0, 1, \dots, n-1\}$
- Heap positions are $\{0, 1, \dots, n-1\}$
- **VtoH** maps vertices to heap positions
- **HtoV** maps heap positions to vertices

- Update node 1 to 35



VtoH


0	1	2	3	4	5	6	7	8
7	3	5	2	4	1	6	0	8

HtoV

0	1	2	3	4	5	6	7	8
7	5	3	1	4	2	6	0	8

Updating values in a min-heap

- Maintain two additional dictionaries
 - Vertices are $\{0, 1, \dots, n-1\}$
 - Heap positions are $\{0, 1, \dots, n-1\}$
 - **VtoH** maps vertices to heap positions
 - **HtoV** maps heap positions to vertices
- Update node 1 to 35
- Update **VtoH** and **HtoV** each time we swap values in the heap



VtoH


0	1	2	3	4	5	6	7	8
7	1	5	2	4	3	6	0	8

HtoV

0	1	2	3	4	5	6	7	8
7	1	3	5	4	2	6	0	8

Dijkstra's algorithm

- Using min-heaps
 - Identifying next vertex to visit is $O(\log n)$
 - Updating distance takes $O(\log n)$ per neighbour
 - Adjacency list — proportionally to degree



VtoH


0	1	2	3	4	5	6	7	8
7	3	5	2	4	1	6	0	8

HtoV

0	1	2	3	4	5	6	7	8
7	5	3	1	4	2	6	0	8

Dijkstra's algorithm

- Using min-heaps
 - Identifying next vertex to visit is $O(\log n)$
 - Updating distance takes $O(\log n)$ per neighbour
 - Adjacency list — proportionally to degree
- Cumulatively
 - $O(n \log n)$ to identify vertices to visit across n iterations
 - $O(m \log n)$ distance updates overall



VtoH


0	1	2	3	4	5	6	7	8
7	3	5	2	4	1	6	0	8

HtoV

0	1	2	3	4	5	6	7	8
7	5	3	1	4	2	6	0	8

Dijkstra's algorithm

- Using min-heaps
 - Identifying next vertex to visit is $O(\log n)$
 - Updating distance takes $O(\log n)$ per neighbour
 - Adjacency list — proportionally to degree
- Cumulatively
 - $O(n \log n)$ to identify vertices to visit across n iterations
 - $O(m \log n)$ distance updates overall
- Overall $O((m + n) \log n)$



VtoH

0	1	2	3	4	5	6	7	8
7	3	5	2	4	1	6	0	8

HtoV

0	1	2	3	4	5	6	7	8
7	5	3	1	4	2	6	0	8

Heap sort

- Start with an unordered list

Heap sort

- Start with an unordered list
- Build a heap — $O(n)$

Heap sort

- Start with an unordered list
- Build a heap — $O(n)$
- Call `delete_max()` n times to extract elements in descending order — $O(n \log n)$

Heap sort

- Start with an unordered list
- Build a heap — $O(n)$
- Call `delete_max()` n times to extract elements in descending order — $O(n \log n)$
- After each `delete_max()`, heap shrinks by 1

Heap sort

- Start with an unordered list
- Build a heap — $O(n)$
- Call `delete_max()` n times to extract elements in descending order — $O(n \log n)$
- After each `delete_max()`, heap shrinks by 1
- Store maximum value at the end of current heap

Heap sort

- Start with an unordered list
- Build a heap — $O(n)$
- Call `delete_max()` n times to extract elements in descending order — $O(n \log n)$
- After each `delete_max()`, heap shrinks by 1
- Store maximum value at the end of current heap
- In place $O(n \log n)$ sort

Summary

- Updating a value in a heap takes $O(\log n)$

Summary

- Updating a value in a heap takes $O(\log n)$
- Need to maintain additional pointers to map values to heap positions and vice versa

Summary

- Updating a value in a heap takes $O(\log n)$
- Need to maintain additional pointers to map values to heap positions and vice versa
- With this extended notion of heap, Dijkstra's algorithm complexity improves from $O(n^2)$ to $O((m + n) \log n)$

Summary

- Updating a value in a heap takes $O(\log n)$
- Need to maintain additional pointers to map values to heap positions and vice versa
- With this extended notion of heap, Dijkstra's algorithm complexity improves from $O(n^2)$ to $O((m + n) \log n)$
- In a similar way, improve Prim's algorithm to $O((m + n) \log n)$

Summary

- Updating a value in a heap takes $O(\log n)$
- Need to maintain additional pointers to map values to heap positions and vice versa
- With this extended notion of heap, Dijkstra's algorithm complexity improves from $O(n^2)$ to $O((m + n) \log n)$
- In a similar way, improve Prim's algorithm to $O((m + n) \log n)$
- Heaps can also be used to sort a list in place in $O(n \log n)$